

Stock_Market_Prediction

August 16, 2018

1 Predicting the Stock Market Prices

1.0.1 Data Description:

In this mission, we'll be working with a csv file containing index prices. Each row in the file contains a **daily record of the price of the S&P500 Index from 1950 to 2015**. The dataset is stored in sphist.csv.

The columns of the dataset are:

Date -- The date of the record.
Open -- The opening price of the day (when trading starts).
High -- The highest trade price during the day.
Low -- The lowest trade price during the day.
Close -- The closing price for the day (when trading is finished).
Volume -- The number of shares traded.
Adj Close -- The daily closing price, adjusted for corporate actions.

We'll be using this dataset to develop a predictive model. We'll train the model with data from 1950-2012, and try to make predictions from 2013-2015.

```
In [50]: from sklearn.linear_model import LinearRegression
import matplotlib.pyplot as plt

# Importing the statistics module
from statistics import mean
from statistics import median

def linearReg(x_train, y_train, x_test, y_test):

    lm = LinearRegression()
    lm.fit(x_train, y_train)

    y_pred = lm.predict(x_test)

    plt.scatter(y_test, y_pred)
    plt.xlabel("Actual Prices: $Y_i$")
    plt.ylabel("Predicted Prices: $\hat{Y}_i$")
    plt.title("Actual Prices vs Predicted prices: $Y_i$ vs $\hat{Y}_i$")
```

```

plt.show()

# Calculating the error
delta_y = y_test - y_pred;
print("Mean Absolute Error (MAE) = " + str(mean(abs(delta_y))))

# Calculating the percentage error
delta_y_percentage = (y_test - y_pred)/ y_test;

print(bold + "Median Absolute Percentage Error (MAPE) = " +
      str(round(median(abs(delta_y_percentage)), 2)*100) + "%" + end)

```

In [54]: `from sklearn.preprocessing import StandardScaler`

```

def trainTestSplit(sp_augmented, predictWeek = False):

    # Train dataset would contain rows with a date less than 2013-01-01
    train = sp_augmented[sp_augmented["Date"] < datetime(year=2013, month=1, day=1)]
    test = sp_augmented[sp_augmented["Date"] >= datetime(year=2013, month=1, day=1)]

    # Separate Train dataset & remove unusable columns for prediction
    # if % change in a week is to be predicted, then response variable
    # should be set to percentage change in closing price over a week.
    if predictWeek:
        y_train = train['closePerChange']
        x_train = train.drop(
            ['closePerChange', 'Close', 'High', 'Low',
             'Open', 'Volume', 'Adj Close', 'Date'], axis=1)
    else:
        y_train = train['Close']
        x_train = train.drop(
            ['Close', 'High', 'Low', 'Open',
             'Volume', 'Adj Close', 'Date'], axis=1)

    # Separate Test dataset & remove unusable columns for prediction
    if predictWeek:
        y_test = test['closePerChange']
        x_test = test.drop(
            ['closePerChange', 'Close', 'High', 'Low',
             'Open', 'Volume', 'Adj Close', 'Date'], axis=1)
    else:
        y_test = test['Close']
        x_test = test.drop(
            ['Close', 'High', 'Low', 'Open',
             'Volume', 'Adj Close', 'Date'], axis=1)

    # Standardisation.

```

```

#     scaler = StandardScaler(copy=False).fit(x_train)
#     x_train = scaler.transform(x_train)

#     scaler = StandardScaler(copy=False)
#     y_train = np.squeeze(scaler.fit_transform(y_train.reshape(-1, 1)))

#     scaler = StandardScaler(copy=False).fit(x_test)
#     x_test = scaler.transform(x_test)

#     scaler = StandardScaler(copy=False)
#     y_test = np.squeeze(scaler.fit_transform(y_test.reshape(-1, 1)))

print("\n\n*** Train Dataset Outcome Variables ***")
print(y_train.head(5))
print("\n\n*** Train Dataset Input Variables ***")
print(x_train.head(5))
print("\n\n*** Test Dataset Outcome Variables ***")
print(y_test.head(5))
print("\n\n*** Test Dataset Input Variables ***")
print(x_test.head(5))

return x_train, y_train, x_test, y_test

```

2 Prediction of Stock Prices: Week Ahead

To predict % change in stock price after a week, we can use features averaging previous 'n' days coupled with % change in stock price after a week, as the response variable, y as training data. In test data also, a 7-day forward shift in the closing price is introduced to compare the % change against prediction.

```

In [56]: # print(sp_augmented.head(5))
# drop the added features for daily prediction
sp_augmented = sp_augmented.drop(['mean_365', 'std_365', 'mean_5by365'], axis=1)

# closePrices.shift(-7) would give the closing price after a week.
closePrices = pd.DataFrame(sp_augmented.Close)
shiftedByWeek = closePrices.shift(-7)

# to find the % change in closing price after a week
closePriceChange = 100 * (shiftedByWeek - closePrices)/closePrices
sp_augmented['closePerChange'] = closePriceChange

# The average price for the past 7 days.
# The standard deviation of the price over the past 7 days
shifted = closePrices.shift(1)
window = shifted.rolling(window=7)
means_7 = window.mean()

```

```

std_7 = window.std()
# sp_augmented['mean_7'] = means_7
# sp_augmented['std_7'] = std_7

window = shifted.rolling(window=30)
means_30 = window.mean()
std_30 = window.std()
sp_augmented['mean_7by30'] = means_7/ means_30
sp_augmented['std_7by30'] = std_7/ std_30

# The ratio between the average price
# for the past 7 days, and the average
# price for the past 14 days.
window = shifted.rolling(window=14)
means_14 = window.mean()
std_14 = window.std()
sp_augmented['mean_7by14'] = means_7/ means_14
sp_augmented['std_7by14'] = std_7/ std_14

# Use the dropna method to remove any rows with NaN values.
# Pass in the axis=0 argument to drop rows.
sp_augmented = sp_augmented.dropna(axis = 0)

print("\n\n*** Dataset After Augmented Values ***")
print(sp_augmented.head(10))

# To split into train and test data to do linear regression
x_train, y_train, x_test, y_test = trainTestSplit(sp_augmented, predictWeek = True)
linearReg(x_train, y_train, x_test, y_test)

```

*** Dataset After Augmented Values ***

	Date	Open	High	Low	Close	Volume	\
16194	1951-08-01	22.510000	22.510000	22.510000	22.510000	1680000.0	
16193	1951-08-02	22.820000	22.820000	22.820000	22.820000	2130000.0	
16192	1951-08-03	22.850000	22.850000	22.850000	22.850000	1570000.0	
16191	1951-08-06	23.010000	23.010000	23.010000	23.010000	1600000.0	
16190	1951-08-07	23.030001	23.030001	23.030001	23.030001	1810000.0	
16189	1951-08-08	22.930000	22.930000	22.930000	22.930000	1410000.0	
16188	1951-08-09	22.840000	22.840000	22.840000	22.840000	1500000.0	
16187	1951-08-10	22.790001	22.790001	22.790001	22.790001	1260000.0	
16186	1951-08-13	22.799999	22.799999	22.799999	22.799999	1320000.0	
16185	1951-08-14	22.700001	22.700001	22.700001	22.700001	1180000.0	

	Adj Close	closePerChange	mean_7by30	std_7by30	mean_7by14	std_7by14
16194	22.510000	1.243896	1.028349	0.378274	1.012455	0.539814

16193	22.820000	-0.087647	1.030264	0.213811	1.012780	0.318141
16192	22.850000	-0.656451	1.031318	0.326064	1.012489	0.468263
16191	23.010000	-0.956102	1.033098	0.327212	1.012252	0.499837
16190	23.030001	-0.694746	1.034323	0.387876	1.012178	0.592194
16189	22.930000	0.043615	1.034844	0.414048	1.011657	0.652498
16188	22.840000	0.394046	1.034237	0.410162	1.010066	0.721764
16187	22.790001	0.175511	1.034788	0.288708	1.009783	0.614055
16186	22.799999	-0.219294	1.033962	0.162876	1.009352	0.398047
16185	22.700001	0.881053	1.030978	0.177855	1.008084	0.426281

*** Train Dataset Outcome Variables ***

16194	1.243896
16193	-0.087647
16192	-0.656451
16191	-0.956102
16190	-0.694746

Name: closePerChange, dtype: float64

*** Train Dataset Input Variables ***

	mean_7by30	std_7by30	mean_7by14	std_7by14
16194	1.028349	0.378274	1.012455	0.539814
16193	1.030264	0.213811	1.012780	0.318141
16192	1.031318	0.326064	1.012489	0.468263
16191	1.033098	0.327212	1.012252	0.499837
16190	1.034323	0.387876	1.012178	0.592194


*** Test Dataset Outcome Variables ***

738	0.658498
737	0.774996
736	0.400281
735	0.734665
734	1.632633

Name: closePerChange, dtype: float64

*** Test Dataset Input Variables ***

	mean_7by30	std_7by30	mean_7by14	std_7by14
738	1.006792	0.707287	0.998234	1.090721
737	1.006252	1.082845	0.998381	1.211434
736	1.007480	1.267217	0.999759	1.292730
735	1.009620	1.403724	1.001383	1.343386
734	1.012163	1.399322	1.003158	1.342442



Mean Absolute Error (MAE) = 1.4682141984586512
Median Absolute Percentage Error (MAPE) = 93.0%

3 MAE vs MAPE Error Metric

MAE error metric is not interpretable, since the value of MAE can range from 0 to infinity. We can't understand how good the model performed. Hence it would be better to compute Percentage error and even better would be to compute median (instead of mean) so that the perturbation caused by outliers could be eliminated.

Median Absolute Percentage Error (MAPE) would be a far more interpretable metric, resilient to outliers.

4 Prediction of Stock Prices: Day Ahead

Percentage error for weekly prediction is unacceptably high. This was expected as stock prediction is an extremely hard problem, to get even a better than random model.

Lets try to predict % change in stock price prior to a day. We can use features, mean and standard deviation of previous 365 days & The ratio between the average price for the past 5 days, and the average price for the past 365 days. The response variable, y , would be the closing price. A 1-day forward shift in the closing price, so that the closing price of the present day (future data) shouldn't be included in prediction.

```

In [55]: import pandas as pd
import numpy as np
from datetime import datetime

# used to format headings
bold = '\033[1m'
end = '\033[0m'

# Read the s&p 500 input data set and sorting based on date.
sp500 = pd.read_csv("sphist.csv", index_col=False)
sp500["Date"] = pd.to_datetime(sp500["Date"])
sp_sorted = sp500.sort_values("Date")

# To do feature engineering - take mean of previous rows
closePrices = pd.DataFrame(sp_sorted.Close)

# Need to shift by 1 coz otherwise it will add up present value also
# Adding present value to the mean is same as giving the
# predicted value (outcome) along with the training data.
# This will hurt prediction model
shifted = closePrices.shift(1)

# The average price for the past 365 days.
# The standard deviation of the price over the past 365 days
window = shifted.rolling(window=365)
means_365 = window.mean()
std_365 = window.std()
sp_sorted['mean_365'] = means_365
sp_sorted['std_365'] = std_365

# The ratio between the average price
# for the past 5 days, and the average
# price for the past 365 days.
window = shifted.rolling(window=5)
means_5 = window.mean()
sp_sorted['mean_5by365'] = means_5/ means_365

# Some of the indicators use 365 days of historical data, and the dataset starts
# on 1950-01-03. Thus, any rows that fall before 1951-01-03 don't have enough
# historical data to compute all the indicators.
sp_augmented = sp_sorted[sp_sorted["Date"] > datetime(year=1951, month=1, day=2)]

# Use the dropna method to remove any rows with NaN values.
# Pass in the axis=0 argument to drop rows.
sp_augmented = sp_augmented.dropna(axis = 0)

print("\n\n*** Dataset After Augmented Values ***")
print(sp_augmented.head(5))

```

```
# To split into train and test data to do linear regression
x_train, y_train, x_test, y_test = trainTestSplit(sp_augmented, predictWeek = False)
linearReg(x_train, y_train, x_test, y_test)
```

*** Dataset After Augmented Values ***

	Date	Open	High	Low	Close	Volume \
16224	1951-06-19	22.020000	22.020000	22.020000	22.020000	1100000.0
16223	1951-06-20	21.910000	21.910000	21.910000	21.910000	1120000.0
16222	1951-06-21	21.780001	21.780001	21.780001	21.780001	1100000.0
16221	1951-06-22	21.549999	21.549999	21.549999	21.549999	1340000.0
16220	1951-06-25	21.290001	21.290001	21.290001	21.290001	2440000.0

	Adj Close	mean_365	std_365	mean_5by365
16224	22.020000	19.447726	1.790253	1.120954
16223	21.910000	19.462411	1.789307	1.125246
16222	21.780001	19.476274	1.788613	1.128142
16221	21.549999	19.489562	1.787659	1.126757
16220	21.290001	19.502082	1.786038	1.121008

*** Train Dataset Outcome Variables ***

```
16224    22.020000
16223    21.910000
16222    21.780001
16221    21.549999
16220    21.290001
Name: Close, dtype: float64
```

*** Train Dataset Input Variables ***


	mean_365	std_365	mean_5by365
16224	19.447726	1.790253	1.120954
16223	19.462411	1.789307	1.125246
16222	19.476274	1.788613	1.128142
16221	19.489562	1.787659	1.126757
16220	19.502082	1.786038	1.121008

*** Test Dataset Outcome Variables ***

```
738    1462.420044
737    1459.369995
736    1466.469971
735    1461.890015
734    1457.150024
Name: Close, dtype: float64
```


*** Test Dataset Input Variables ***

	mean_365	std_365	mean_5by365
738	1327.534055	90.463948	1.068629
737	1327.908247	90.738976	1.073714
736	1328.224877	90.995857	1.079412
735	1328.557617	91.279049	1.086423
734	1328.898603	91.544368	1.095093



Mean Absolute Error (MAE) = 98.94921147104225

Median Absolute Percentage Error (MAPE) = 6.0%

MAPE of 6.0% is a much better prediction result than the previous weekly prediction attempt.

5 Conclusion

- 1) **In daily forecast, Actual and Predicted prices are almost linear. Hence the daily prediction model is working fine**, though the error can be further reduced with a better model such as randomforest or using feature engineering techniques such as previous volume, highest/lowest price in the past year etc.

- 2) For a financial company, another way to reframe this problem would be **to perceive the problem as a classification problem, instead of regression problem**. If we can predict, whether the price of a particular stock would go up or down, on next day or a period of time then such a system is very useful.
- 3) There was a Kaggle competition on similar lines and all of the good solutions just predicted 0% change in price most of the times.

6 Potential Improvements

- a) Learn the domain more and engineer very domain specific features.
- b) Implement this paper for a Deep Learning based momentum trading strategy: Applying Deep Learning To Enhance Momentum Trading Strategies In Stocks by Lawrence Takeuchi

<http://cs229.stanford.edu/proj2013/TakeuchiLee-ApplyingDeepLearningToEnhanceMomentumTradingStr>